

CSCI 2400 – Computer Systems



Instructor:	Nicholas V. Lewchenko (Nick)	nile1033@colorado.edu
TA:	Mayank Joshi	Mayank.Joshi@colorado.edu
TA:	Insoo Lee	Insoo.Lee@colorado.edu
TA:	Austin Pahl	Austin.Pahl@colorado.edu
GSS:	Amatullah Sethjiwala	Amatullah.Sethjiwala@colorado.edu
GSS:	Shruthi Sridharan	Shruthi.Sridharan@colorado.edu

First things first

- **This is a challenging class!**
- **A full semester in 8 weeks.**
- **Start assignments as soon as they are posted!**
- **First big lab is due next Monday (2020-06-08)**

Recitations

■ Recitations:

- 301-REC: Wed, 9:15am-10:35am
- 302-REC: Wed, 11:00am-12:00pm
- 303-REC: Wed, 9:15am-10:35am

■ Check mycuinfo for your section, Moodle for Zoom link

- Can't find your section? You're probably in 303-REC.

Grading: Labs

- **5 Labs, total 60% of course grade**
 - Interview graded, lots of work!
 - First posted later today, due next Monday night
- **Turned in via GitHub Classroom**

In general, things due at 10:00pm

Nothing accepted late!*

Grading: Exams

- **2 Exams, total 32% of course grade**
 - Conducted through Moodle
 - 24-hour window, single 80-minute attempt
- **“Open Book”, using:**
 - Your notes
 - The lecture slides/video
 - Textbook
 - Simple calculator

Grading: Self-Study Quizzes

- **Weekly quizzes, total 8% of course grade**
 - These cover textbook readings and lecture
 - *Computer Systems, A Programmer's Perspective* (3rd Edition) by Bryant and O'Halloran
 - Allow multiple attempts
 - Released Monday morning, due Sunday night
 - This week, four quizzes covering Chapters 1,2

The Many Online Services

■ Moodle

- Exams, quizzes, lecture recordings

■ Piazza

- Weekly overviews, asynchronous office hours

■ Zoom

■ `coding.csel.io`

- Environment for working on labs
- Tutorial in this lecture, more in recitation

■ GitHub Classroom

- Turn-in system for labs, uses Git!
- Invite later today, tutorial in recitation

Important Dates

- **Exam 1:** Thursday, 2020-07-02
- **Exam 2:** Friday, 2020-07-24 (last day of class)
- **Last day to drop:** Monday, 2020-06-08 (1 week from now!)



Course Overview

These slides adapted from materials provided by the

Overview

- Course theme
- Five realities

Course Theme:

Abstraction Is Good But Don't Forget Reality


- **Most CS and CE courses emphasize abstraction**
 - Abstract data types
 - Asymptotic analysis
- **These abstractions have limits**
 - Especially in the presence of bugs
 - Need to understand details of underlying implementations
- **Useful outcomes**
 - Become more effective programmers
 - Able to find and eliminate bugs efficiently
 - Able to understand and tune for program performance
 - Prepare for later “systems” classes in CS & ECE
 - Compilers, Operating Systems, Networks, Computer Architecture, Embedded Systems, Storage Systems, etc.

Great Reality #1:

`int` is not Integers, `float` is not Reals

■ Example 1: Is $x^2 \geq 0$?

■ Float's: Yes!



■ Int's:

- $40000 * 40000$ 1600000000
- $50000 * 50000$??

■ Example 2: Is $(x + y) + z = x + (y + z)$?

■ Unsigned & Signed Int's: Yes!

■ Float's:

- $(1e20 + -1e20) + 3.14$ 3.14
- $1e20 + (-1e20 + 3.14)$??

Computer Arithmetic

- **Does not generate random values**
 - Arithmetic operations have important mathematical properties
- **Cannot assume all the *usual* properties**
 - Necessary due to finiteness of data representations
 - Integer operations satisfy “ring” properties
 - Commutativity, Associativity, Distributivity
 - Floating point operations satisfy “ordering” properties
 - Monotonicity, values of signs

Great Reality #2:

You've Got to Know Assembly

- **Chances are, you'll never write programs in assembly**
 - Compilers are much better & more patient than you are
- **But: Understanding assembly is key to machine-level execution model**
 - Behavior of programs in presence of bugs
 - High-level language models break down
 - Tuning program performance
 - Understand optimizations done / not done by the compiler
 - Understanding sources of program inefficiency
 - Implementing system software
 - Compiler has machine code as target
 - Operating systems must manage process state
 - Creating / fighting malware

Great Reality #3: Memory Matters

Random Access Memory Is an Unphysical Abstraction

- **Memory is not unbounded**
 - It must be allocated and managed
 - Many applications are memory dominated
- **Memory referencing bugs especially pernicious**
 - Effects are distant in both time and space
- **Memory performance is not uniform**
 - Cache and virtual memory effects can greatly affect program performance
 - Adapting program to characteristics of memory system can lead to major speed improvements

Memory Referencing Bug Example

```
typedef struct {
    int a[2];
    double d;
} struct_t;

double fun(int i) {
    volatile struct_t s;
    s.d = 3.14;
    s.a[i] = 1073741824; /* Possibly out of bounds */
    return s.d;
}
```

fun(0)	3.14
fun(1)	3.14
fun(2)	3.1399998664856
fun(3)	2.000000061035156
fun(4)	3.14
fun(6)	Segmentation fault

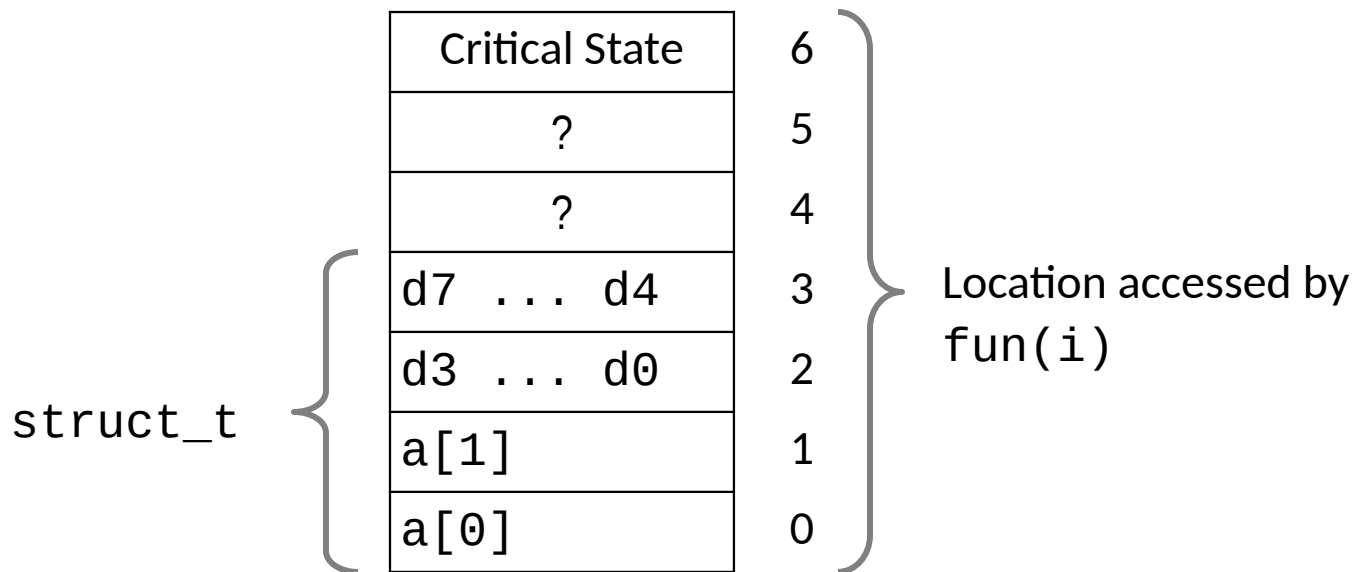
- Result is system specific

Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;
```

fun(0)	3.14
fun(1)	3.14
fun(2)	3.1399998664856
fun(3)	2.000000061035156
fun(4)	3.14
fun(6)	Segmentation fault

Explanation:



Memory Referencing Errors

- **C and C++ do not provide any memory protection**

- Out of bounds array references
- Invalid pointer values
- Abuses of malloc/free

- **Can lead to nasty bugs**

- Whether or not bug has any effect depends on system and compiler
- Action at a distance
 - Corrupted object logically unrelated to one being accessed
 - Effect of bug may be first observed long after it is generated

- **How can I deal with this?**

- Program in Java, Ruby, Python, ML, ...
- Understand what possible interactions may occur
- Use or develop tools to detect referencing errors (e.g. Valgrind)

Great Reality #4: There's more to performance than asymptotic complexity

- **Constant factors matter too!**
- **And even exact op count does not predict performance**
 - Easily see 10:1 performance range depending on how code written
 - Must optimize at multiple levels: algorithm, data representations, procedures, and loops
- **Must understand system to optimize performance**
 - How programs compiled and executed
 - How to measure program performance and identify bottlenecks
 - How to improve performance without destroying code modularity and generality

Memory System Performance Example

```
void copyij(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (i = 0; i < 2048; i++)
        for (j = 0; j < 2048; j++)
            dst[i][j] = src[i][j];
}
```

4.3ms


2.0 GHz Intel Core i7 Haswell

```
void copyji(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (j = 0; j < 2048; j++)
        for (i = 0; i < 2048; i++)
            dst[i][j] = src[i][j];
}
```

81.8ms

- Hierarchical memory organization
- Performance depends on access patterns
 - Including how step through multi-dimensional array

Why The Performance Differs



Great Reality #5:

Computers do more than execute programs

- **They need to get data in and out**
 - I/O system critical to program reliability and performance
- **They communicate with each other over networks**
 - Many system-level issues arise in presence of network
 - Concurrent operations by autonomous processes
 - Coping with unreliable media
 - Cross platform compatibility
 - Complex performance issues

Welcome and Enjoy!