

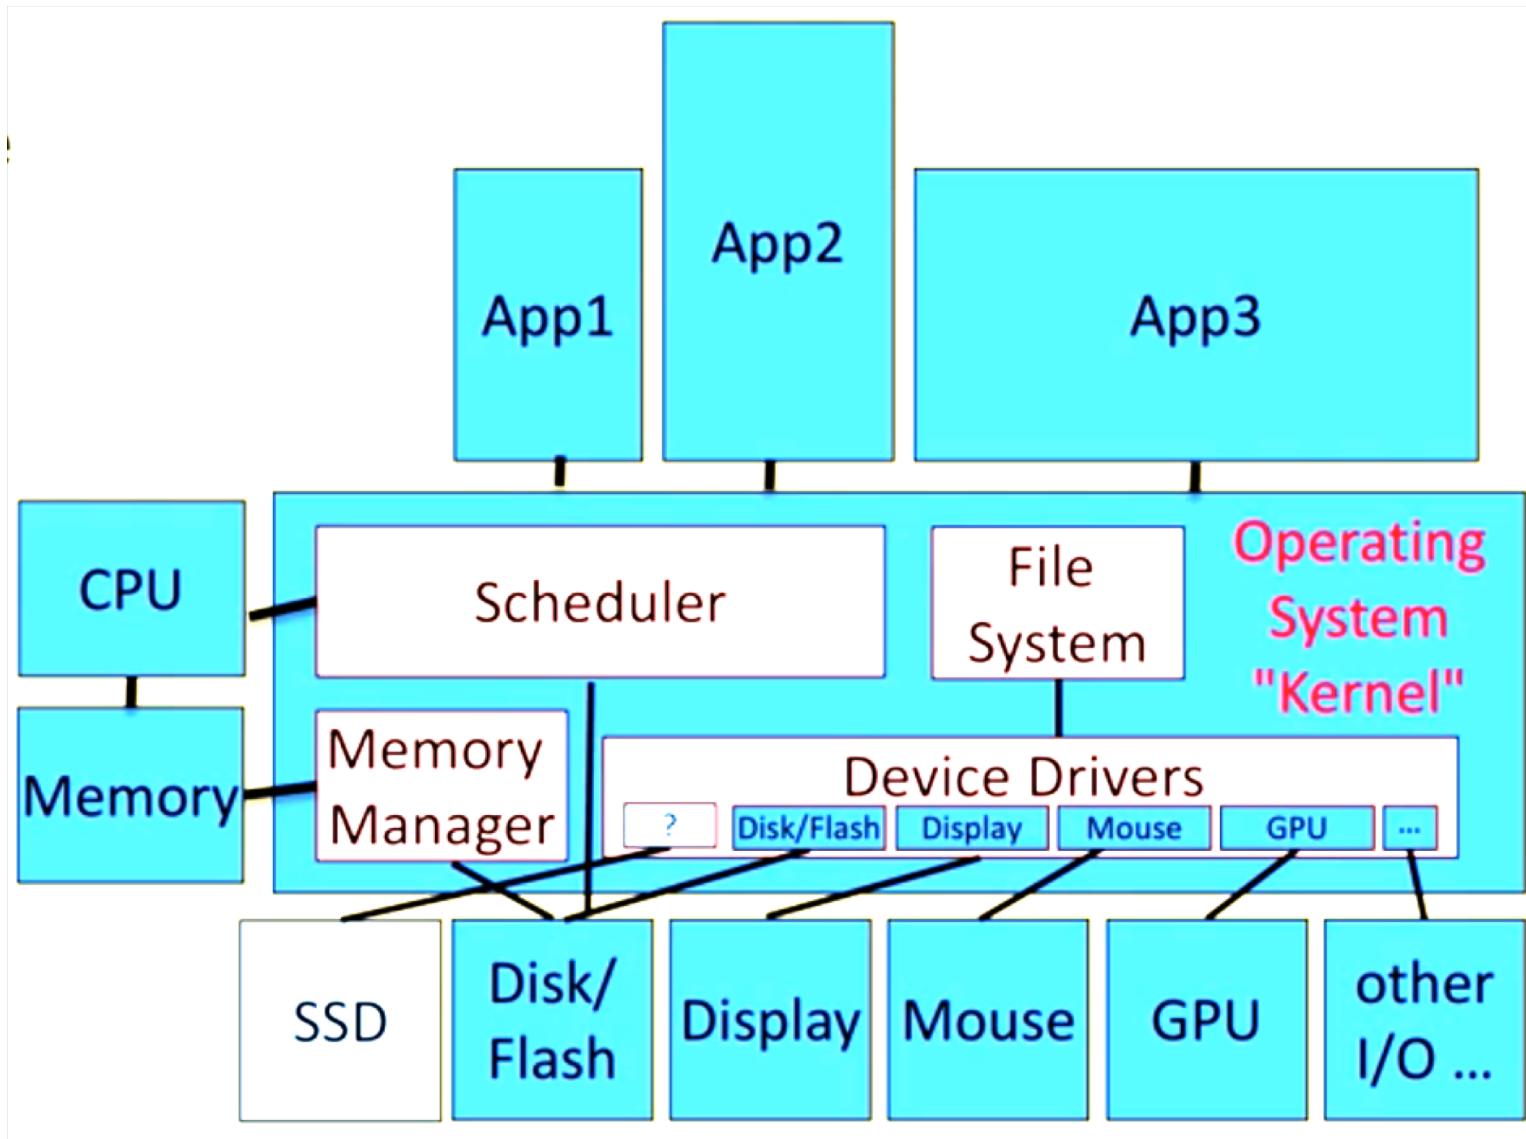


# Midterm Review



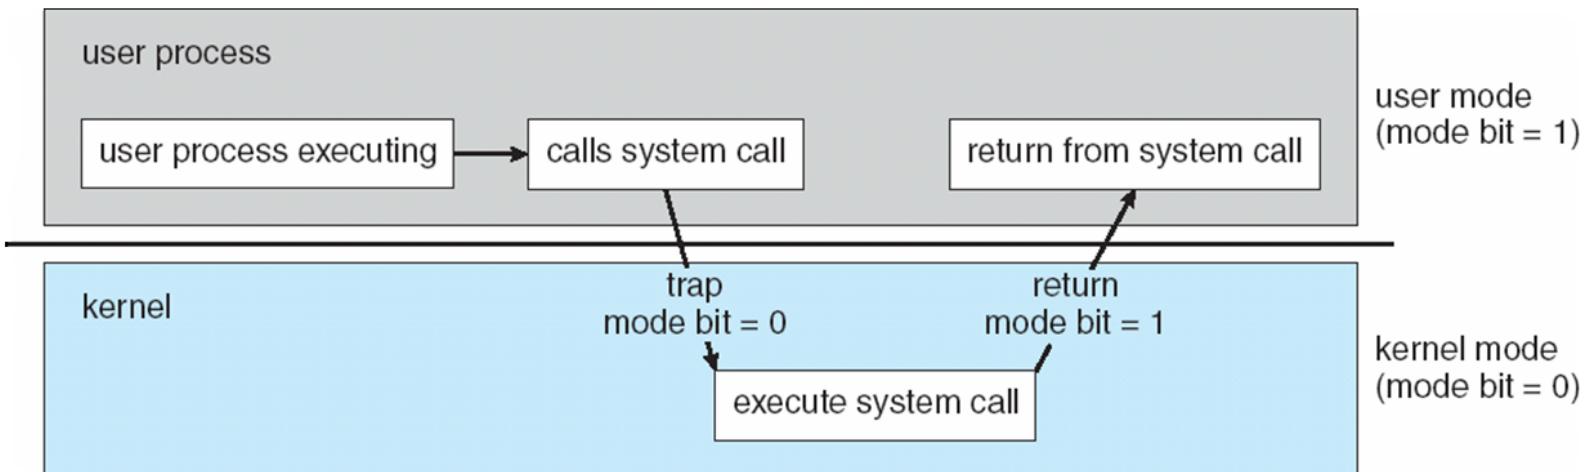
# Overview OS

# Operating Systems (OS)



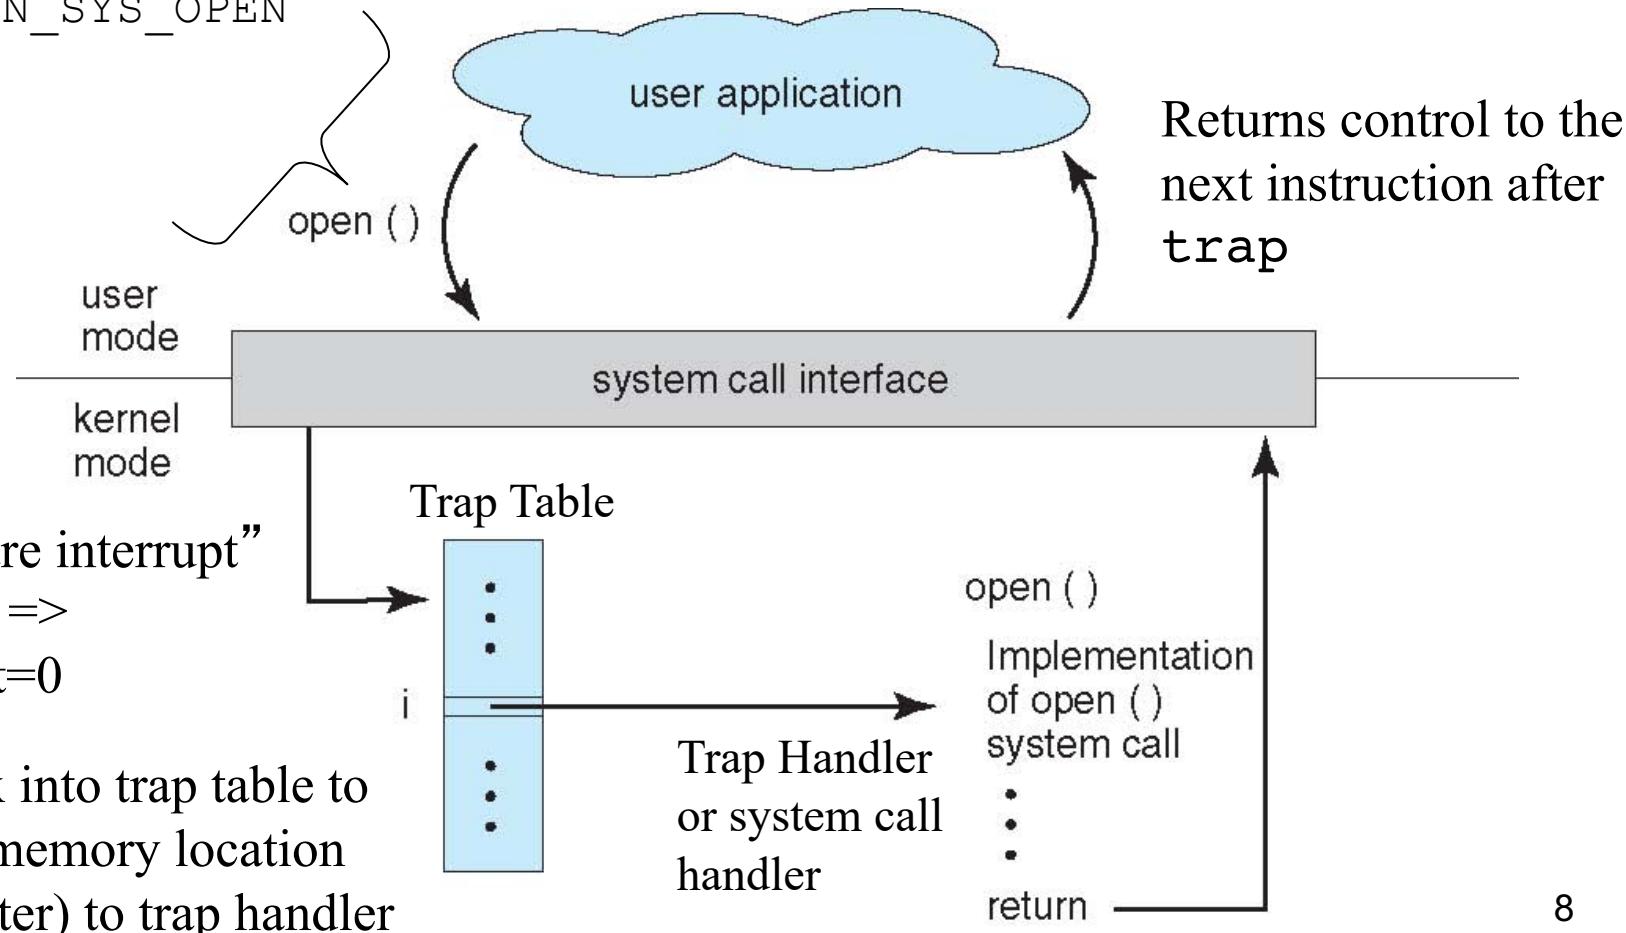
# System Calls

- The `trap` instruction is used to switch from user to supervisor mode, thereby entering the OS
  - `trap` sets the mode bit to 0
  - On x86, use `INT` assembly instruction (more recently `SYSCALL/SYSENTER`)
  - mode bit set back to 1 on return
- Any instruction that invokes `trap` is called a *system call*
  - There are many different classes of system calls

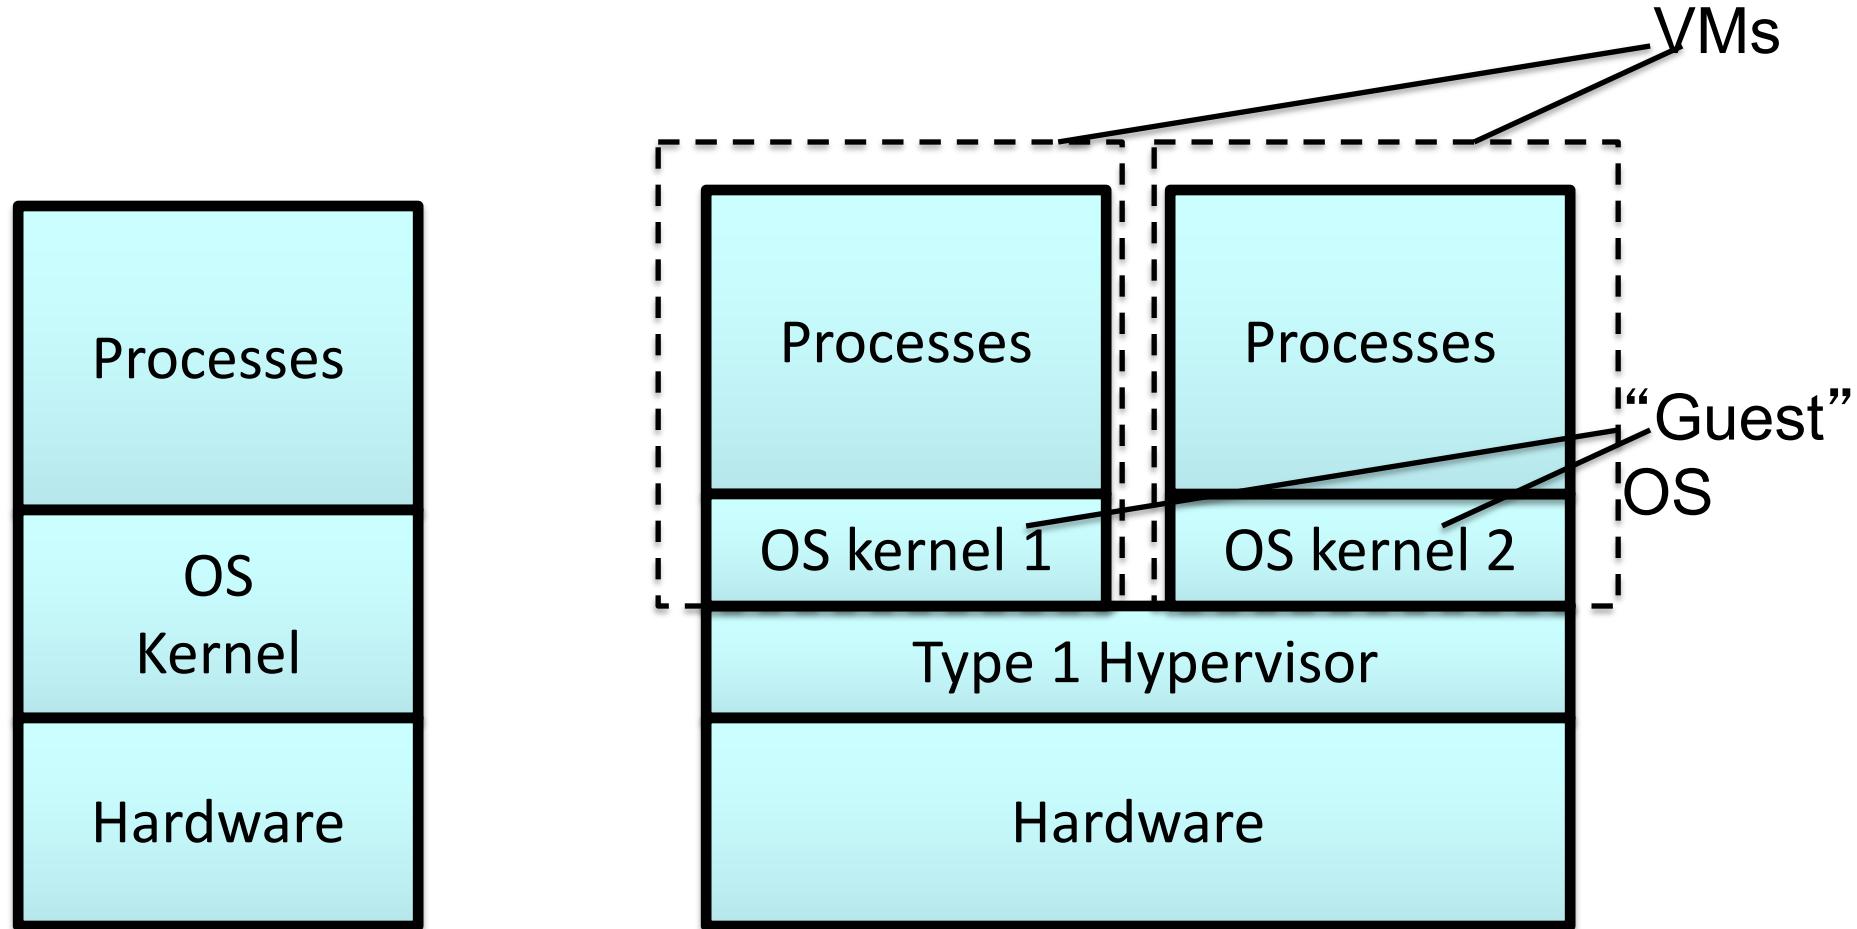


# API – System Calls – OS Relationship

```
open() {  
...  
trap N_SYS_OPEN  
...  
}
```



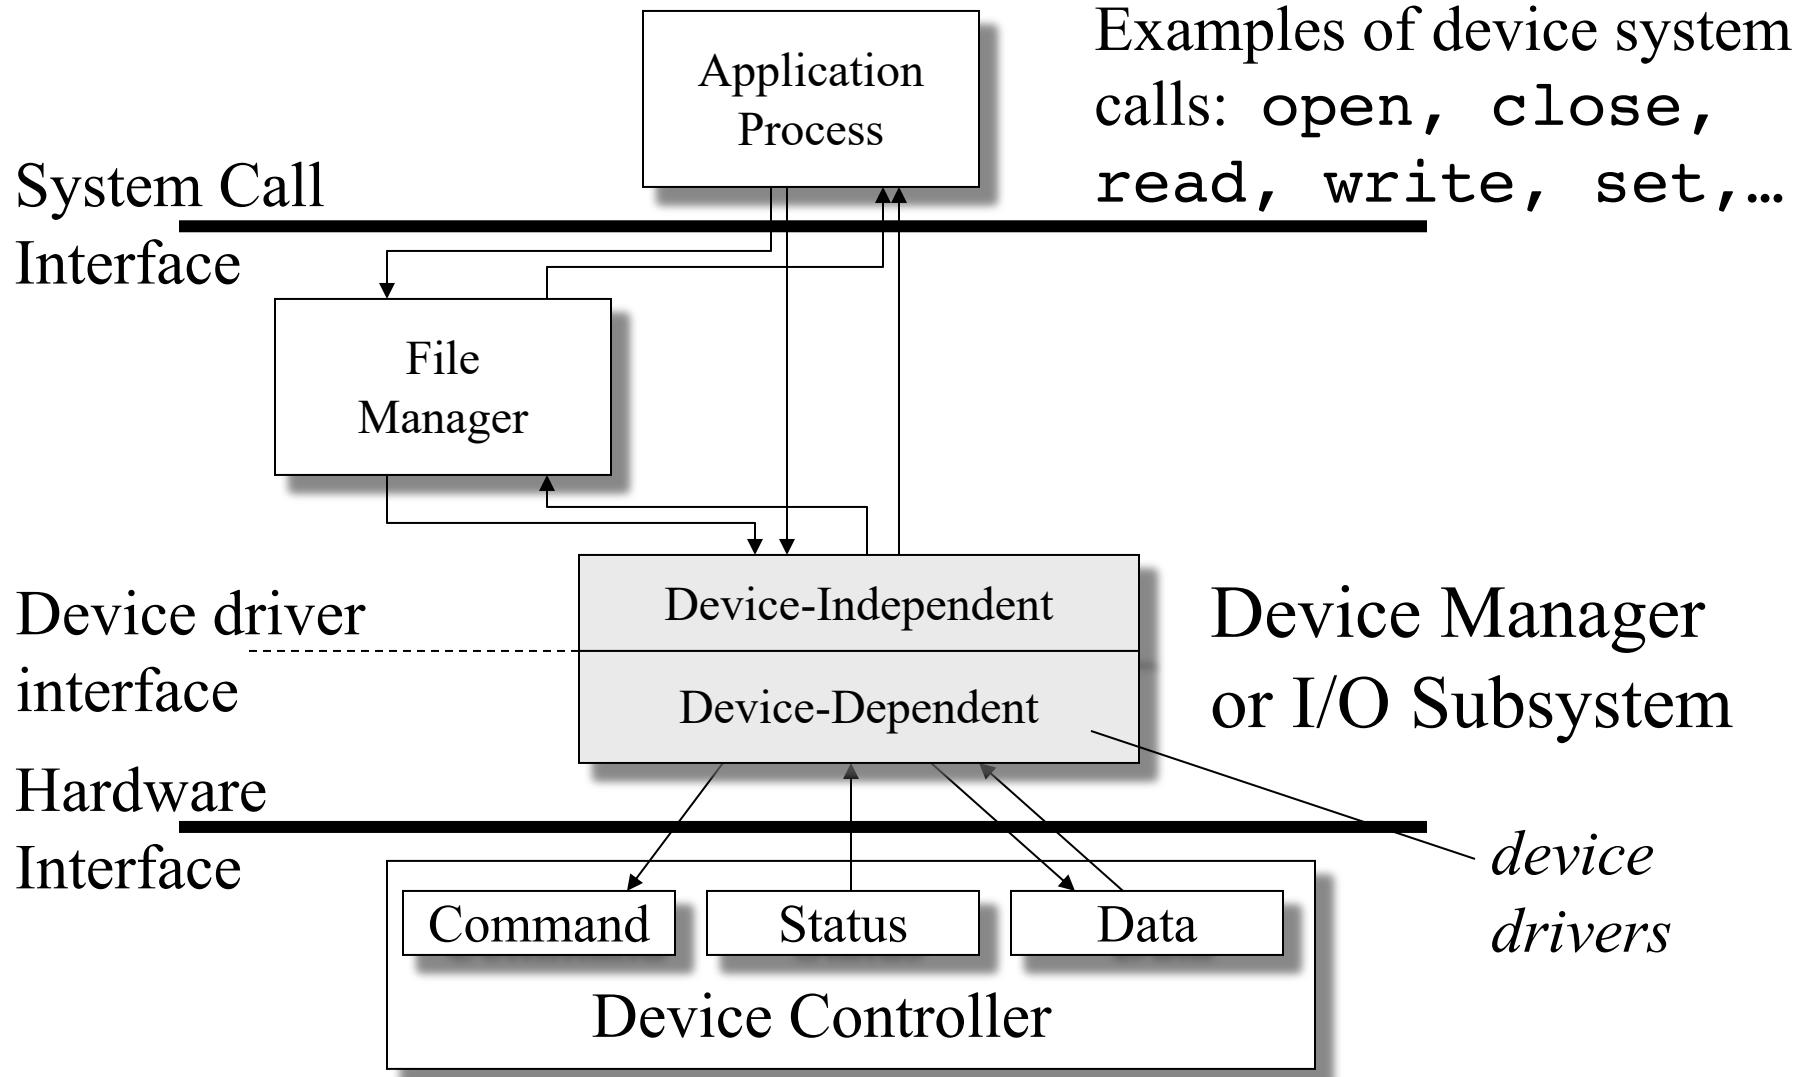
# Virtual Machines



Traditional OS

A *Type 1 Hypervisor* provides a virtualization layer for guest OSs and resides just above the hardware.

# Device Management Organization



# Loadable Kernel Modules (LKM)

- LKM is an object file that contains code to extend a running kernel
- LKMs can be loaded and unloaded from kernel on demand at runtime
- LKMs offer an easy way to extend the functionality of the kernel without having to rebuild or recompile the kernel again
- LKMs are a simple and efficient way to create programs that reside in the kernel and run in privileged mode
- Most of the drivers are written as LKMs

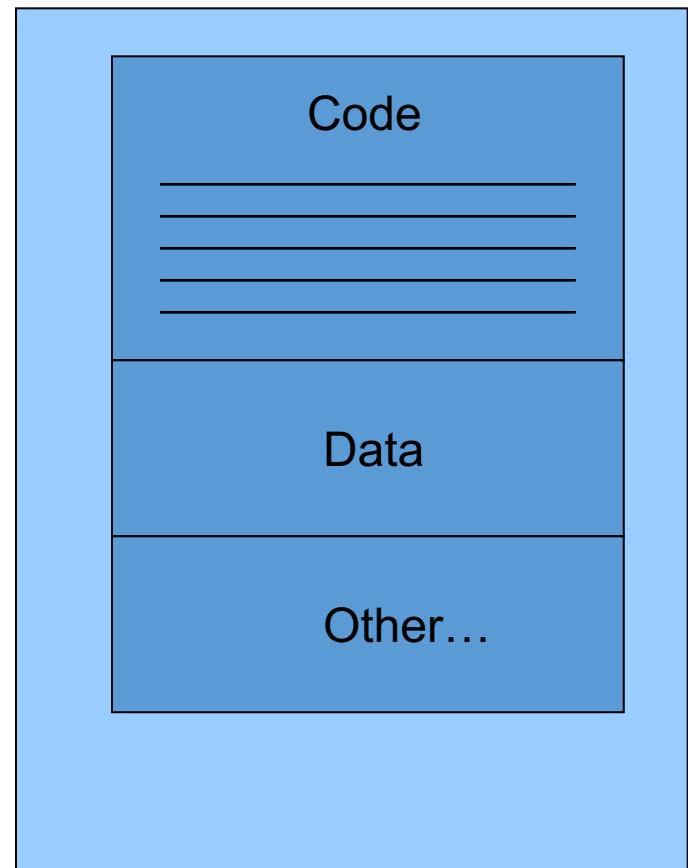


# Process

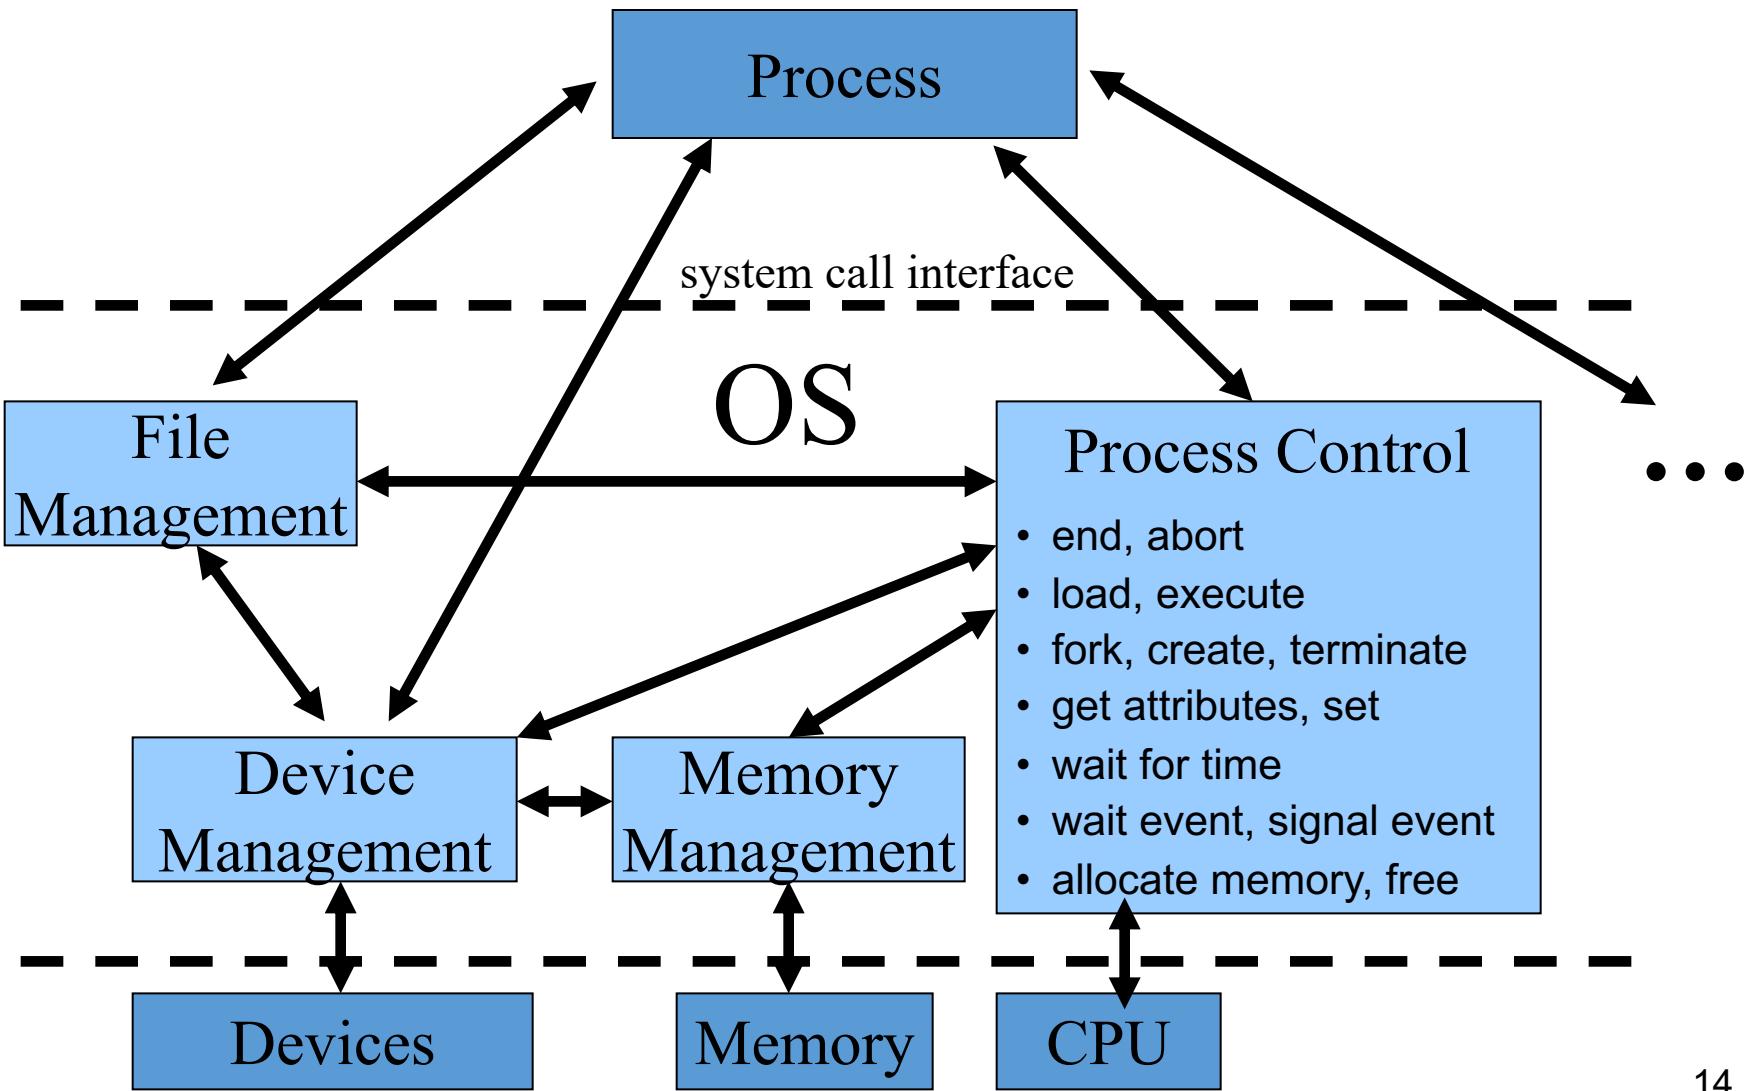
# Process

- A software program consists of a sequence of code instructions and data stored on disk
- A program is a passive entity
- A process is a program actively executing from main memory within its own address space

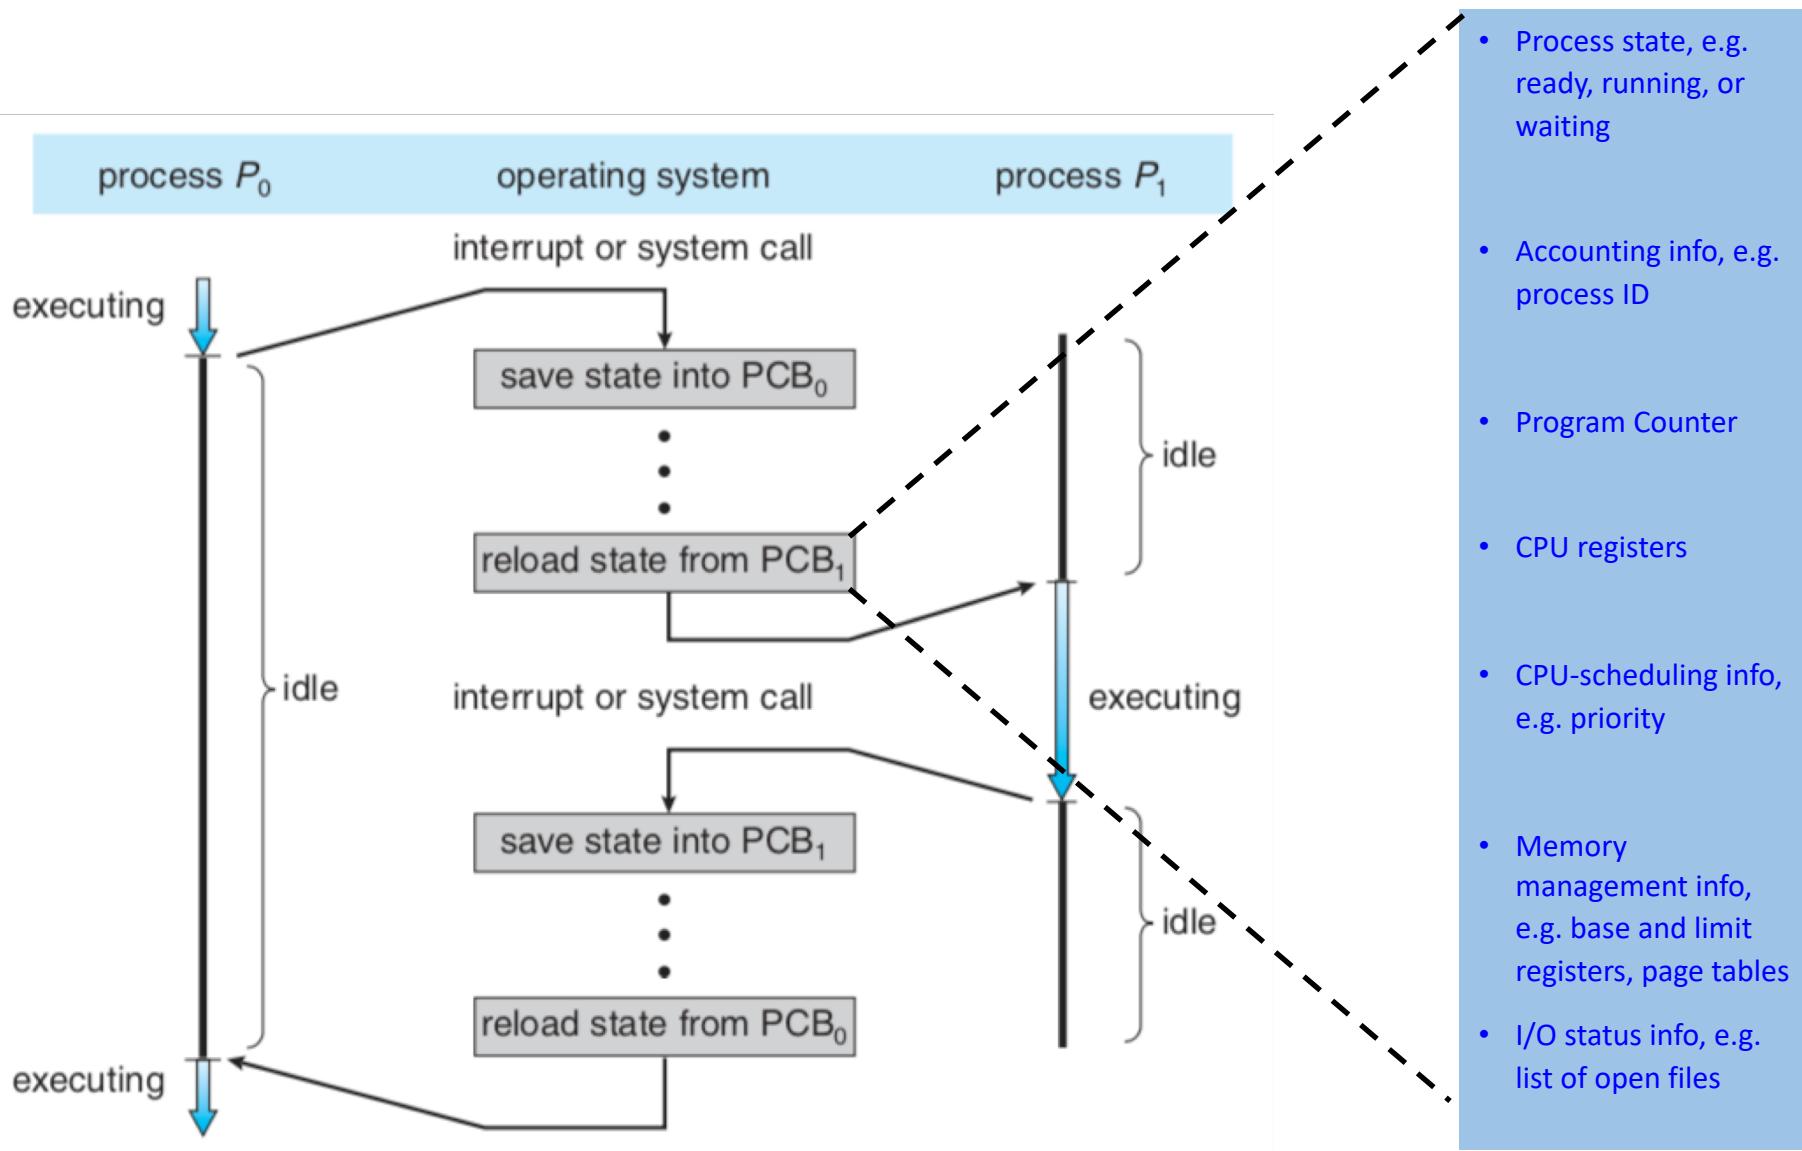
Program P1



# Process Management



# Context Switching

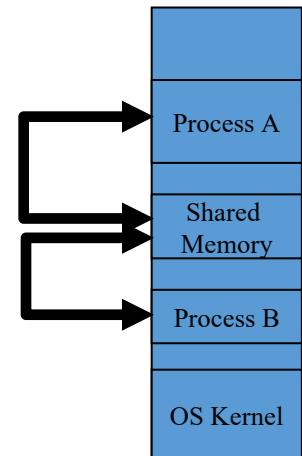
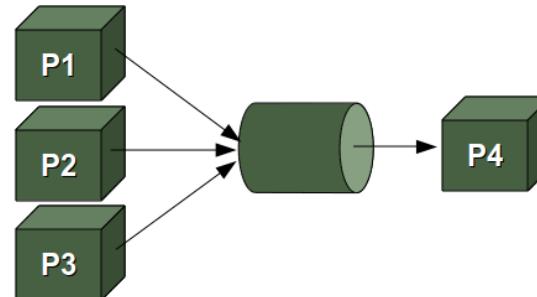




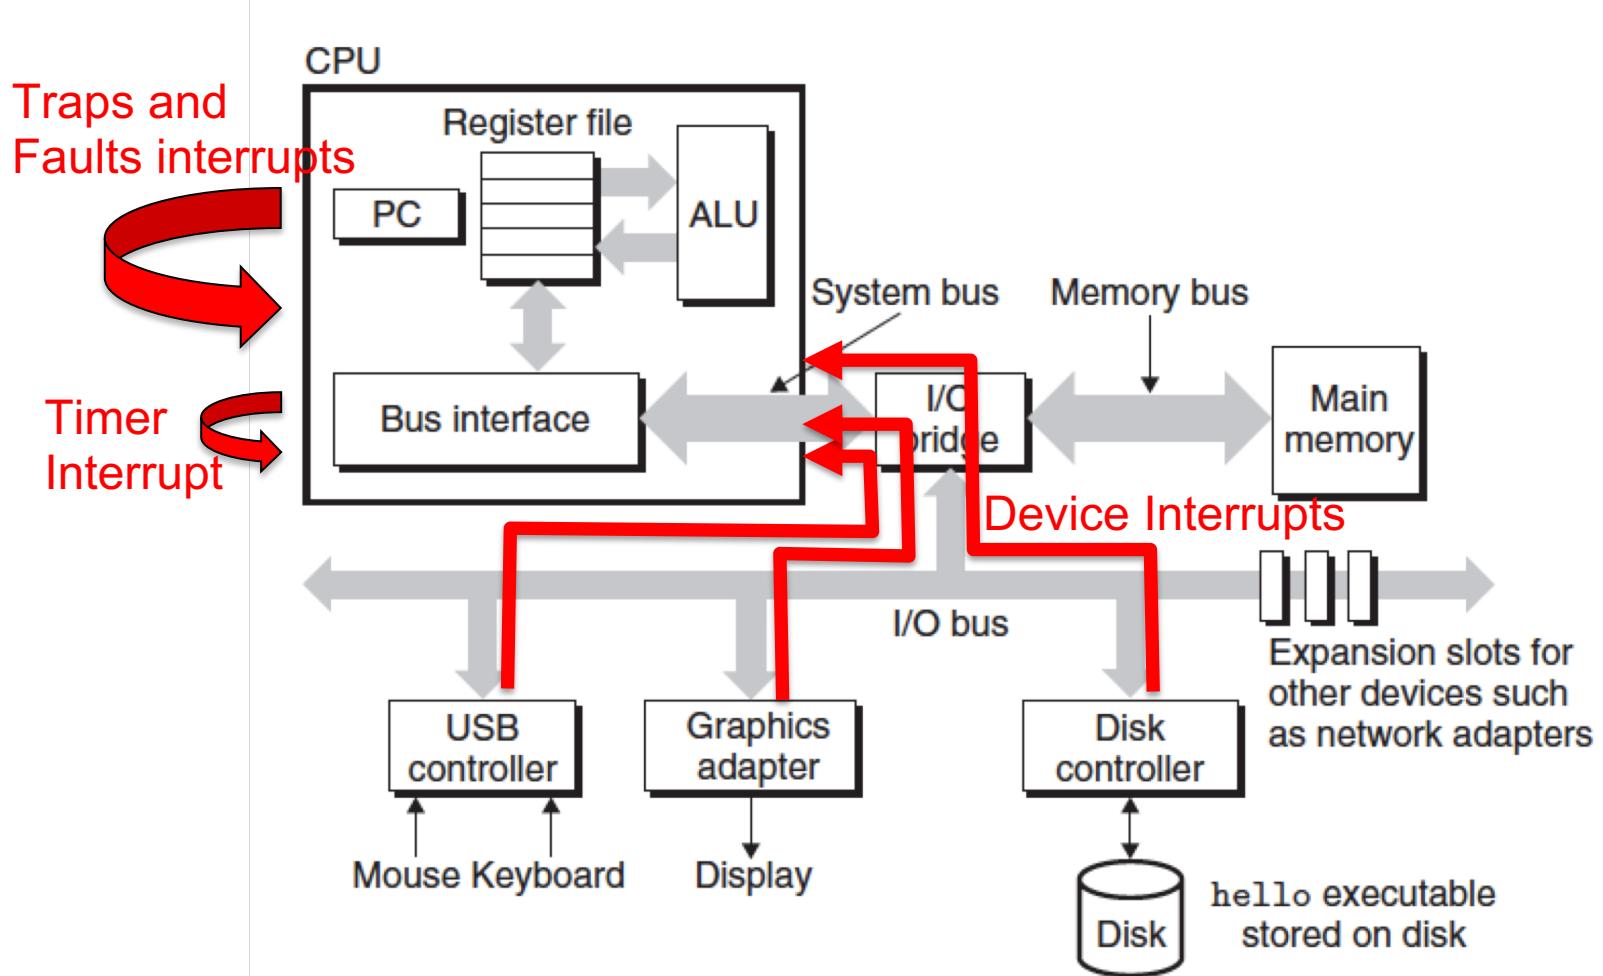
# Interprocess Communication (IPC)

# Interprocess Communication (IPC)

- Signals / Interrupts
  - Notifying that an event has occurred
- Message Passing
  - Pipes
  - Sockets
- Shared Memory
  - Race conditions
  - Synchronization
- Remote Procedure Calls



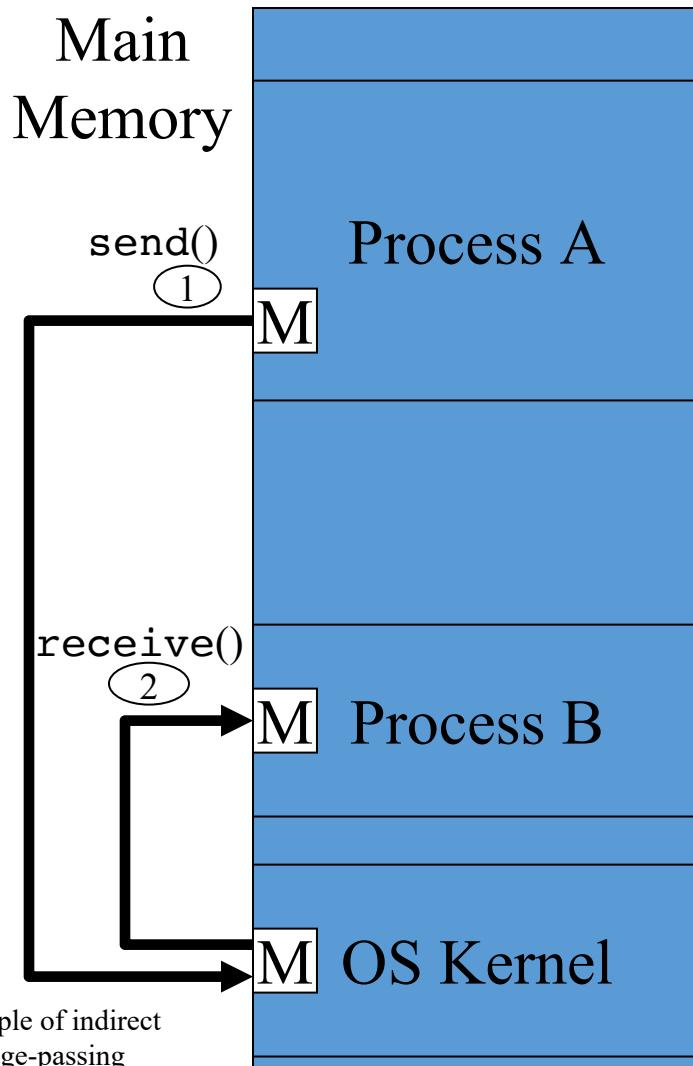
# Linux Signals & Interrupts



# Signals

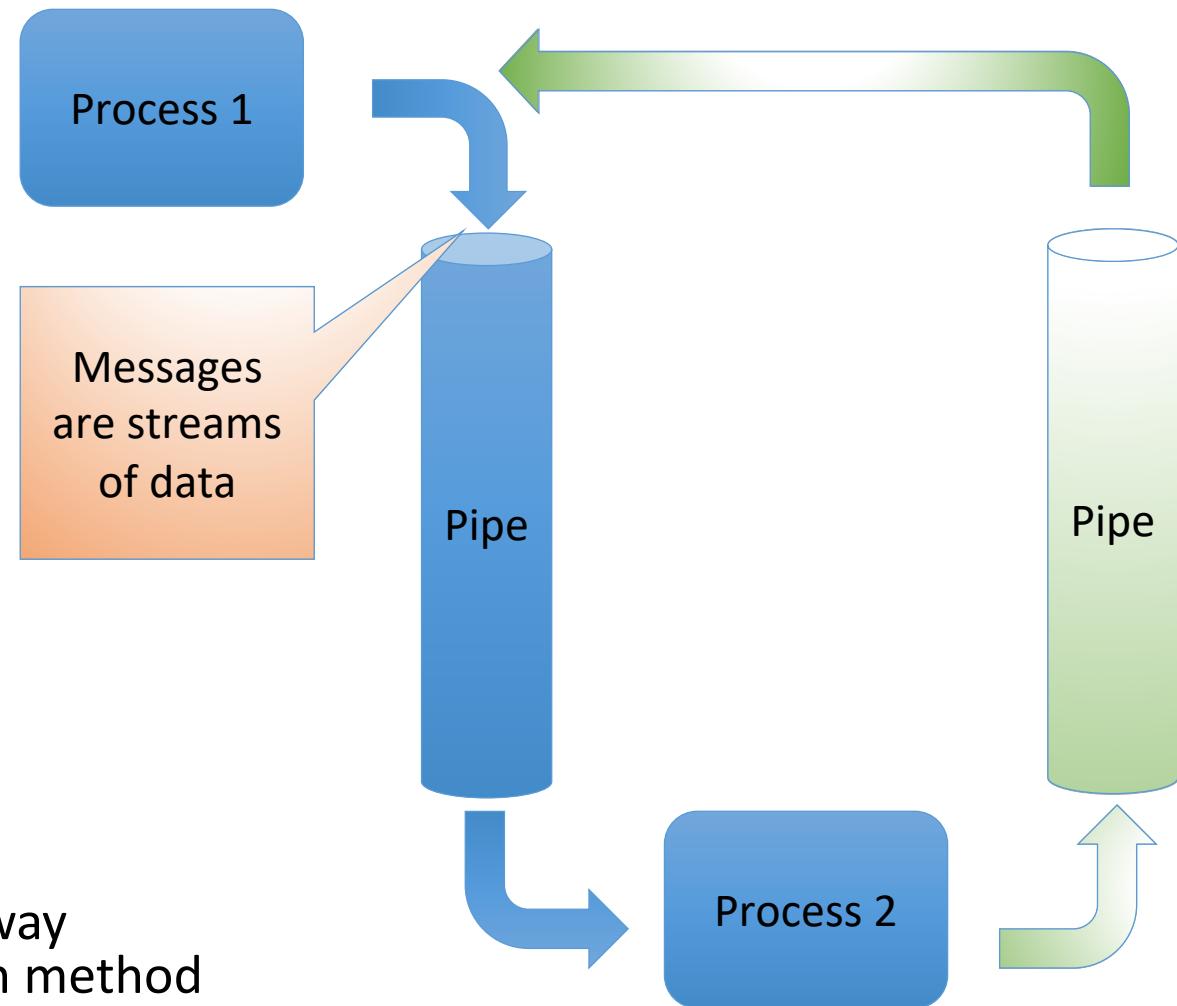
- Kernel-to-Process
- Process-to-Process
- Signals' issues
  - Blocking vs Non-blocking system calls
  - Synchronization vs Asynchronization
  - Race conditions

# IPC Message Passing

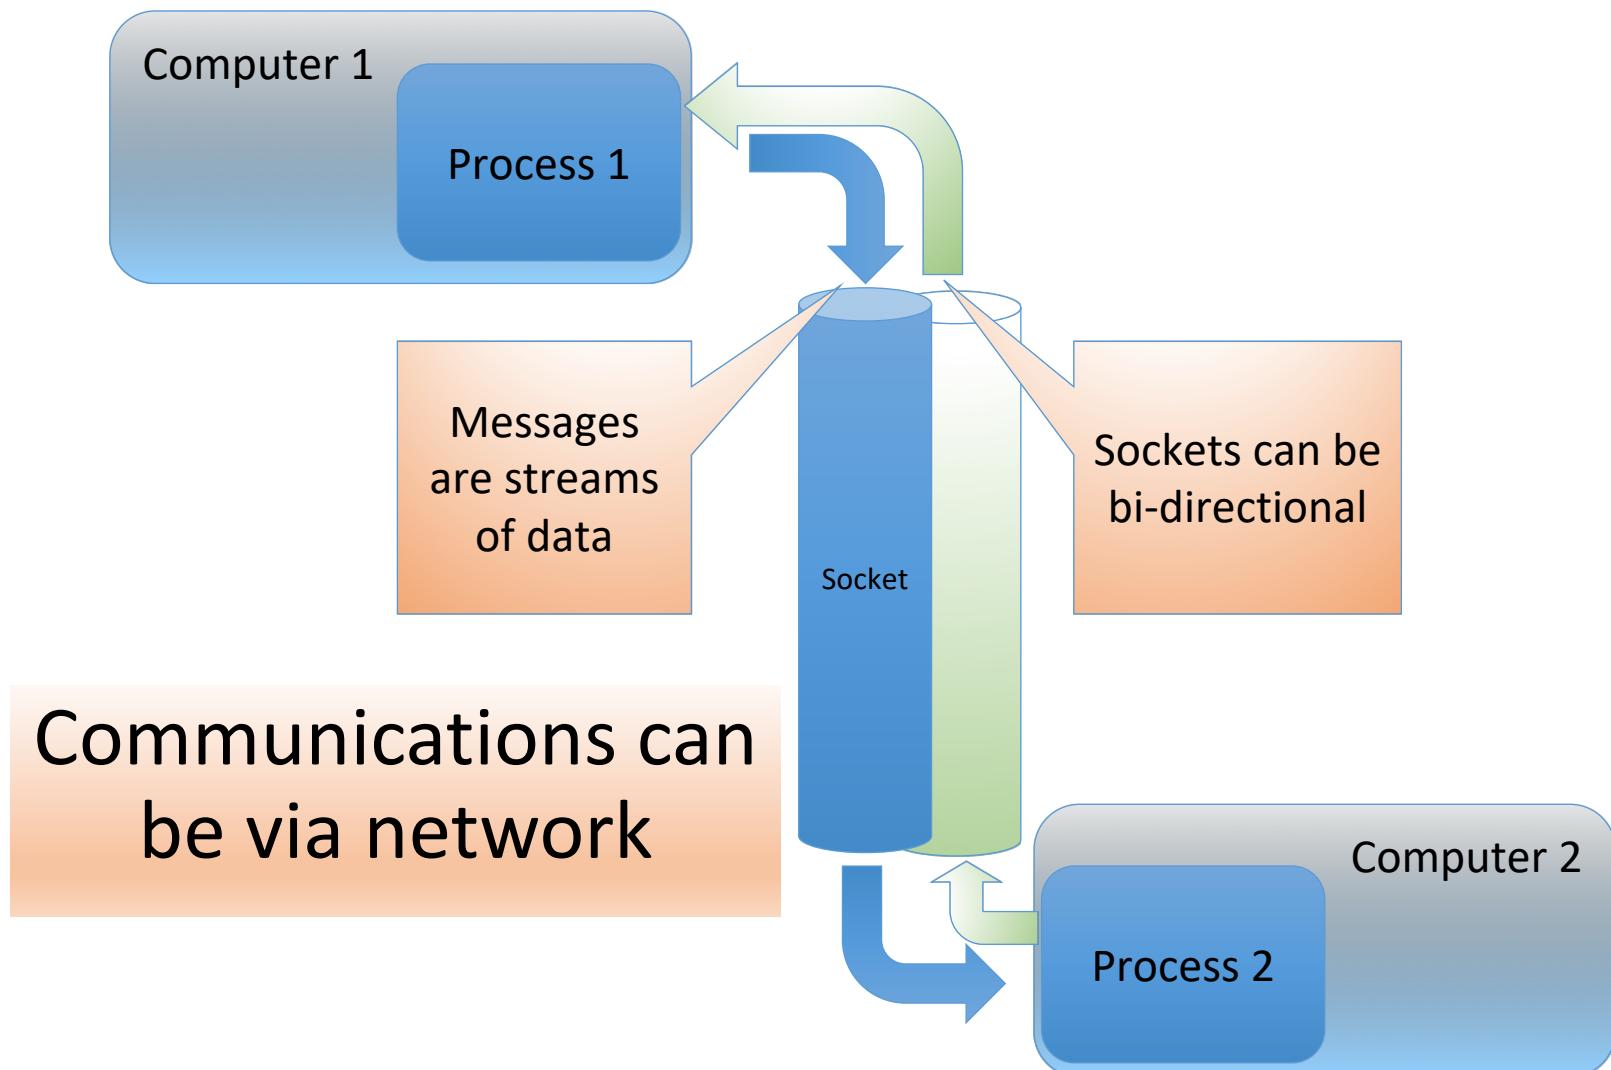


- `send()` and `receive()` can be blocking/synchronous or non-blocking/asynchronous
- Used to pass small messages
- Advantage: OS handles synchronization
- Disadvantage: Slow
  - OS is involved in each IPC operation for control signaling and possibly data as well
- Message Passing IPC types:
  - Pipes
  - UNIX-domain sockets
  - Internet domain sockets
  - message queues
  - remote procedure calls (RPC)

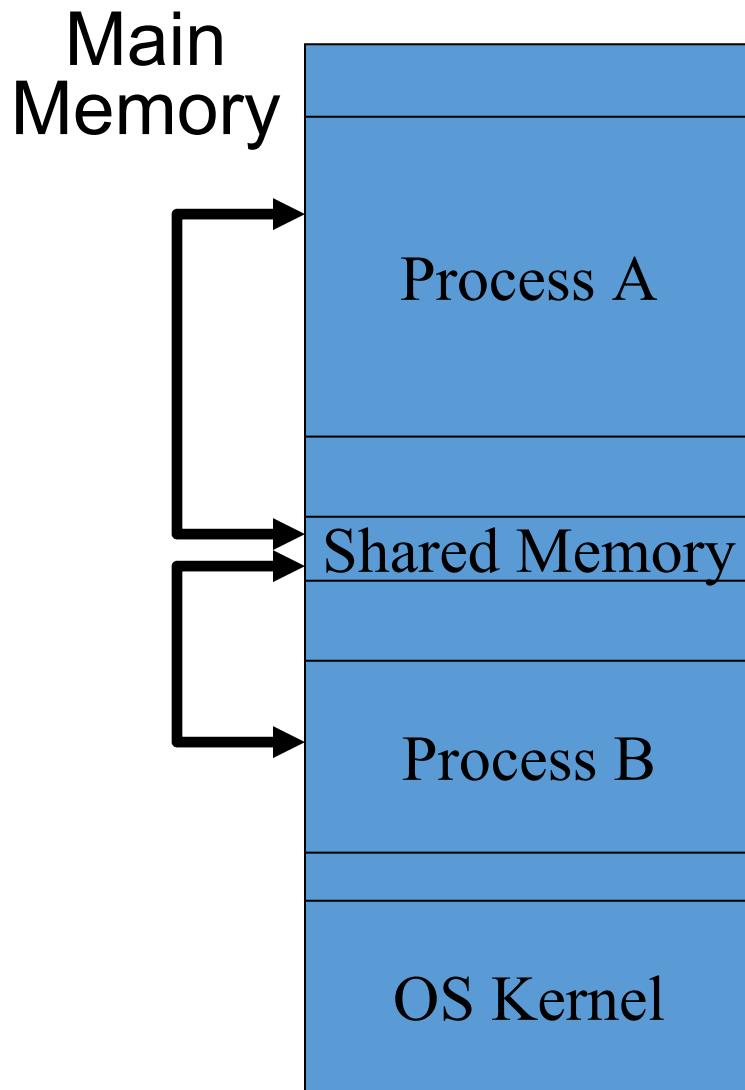
# Message Passing via Pipes



# Message Passing via Sockets



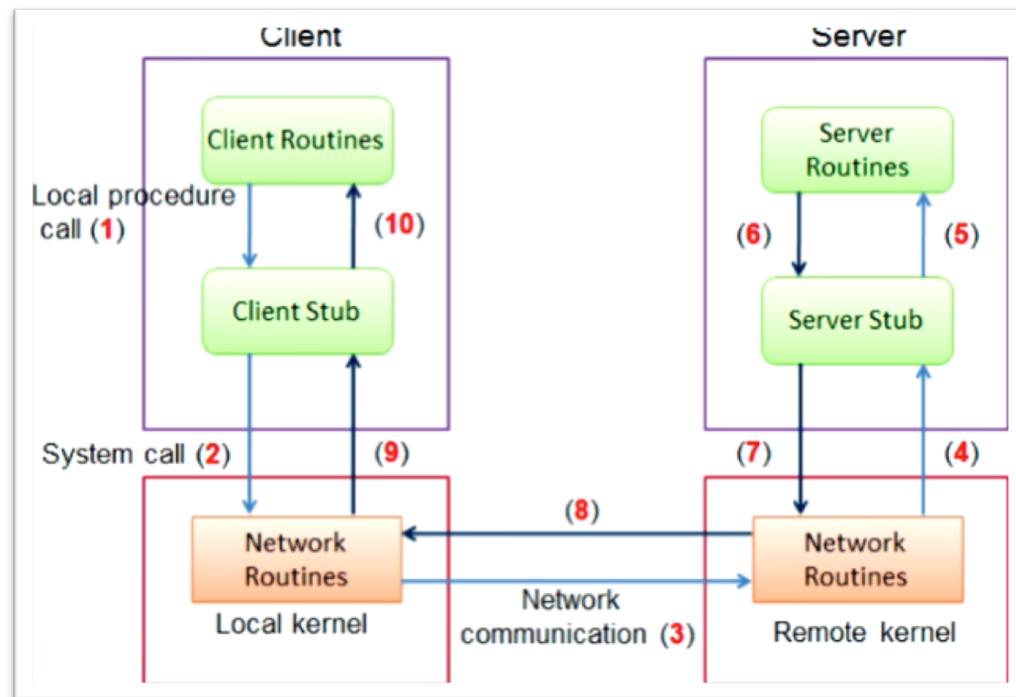
# Shared Memory



- OS provides mechanisms for creation of a shared memory buffer between processes (both processes have address mapped to their process)
- Applies to processes on the same machine
- Problem: shared access introduces complexity
  - need to synchronize access

# Remote Procedure Calls (RPC)

- Client makes a call to a function and passes parameters
- The client has linked a stub for the function. This stub will marshal (packetize) the data and send it to a remote server
- The network transfers the information to a server
- A service listening to the network receives a request
- The information is unmarshalled (unpacked) and the server's function is called.
- The results are returned to be marshalled into a packet
- The network transfers the packet is sent back to the client
- TCP/IP used to transmit packet

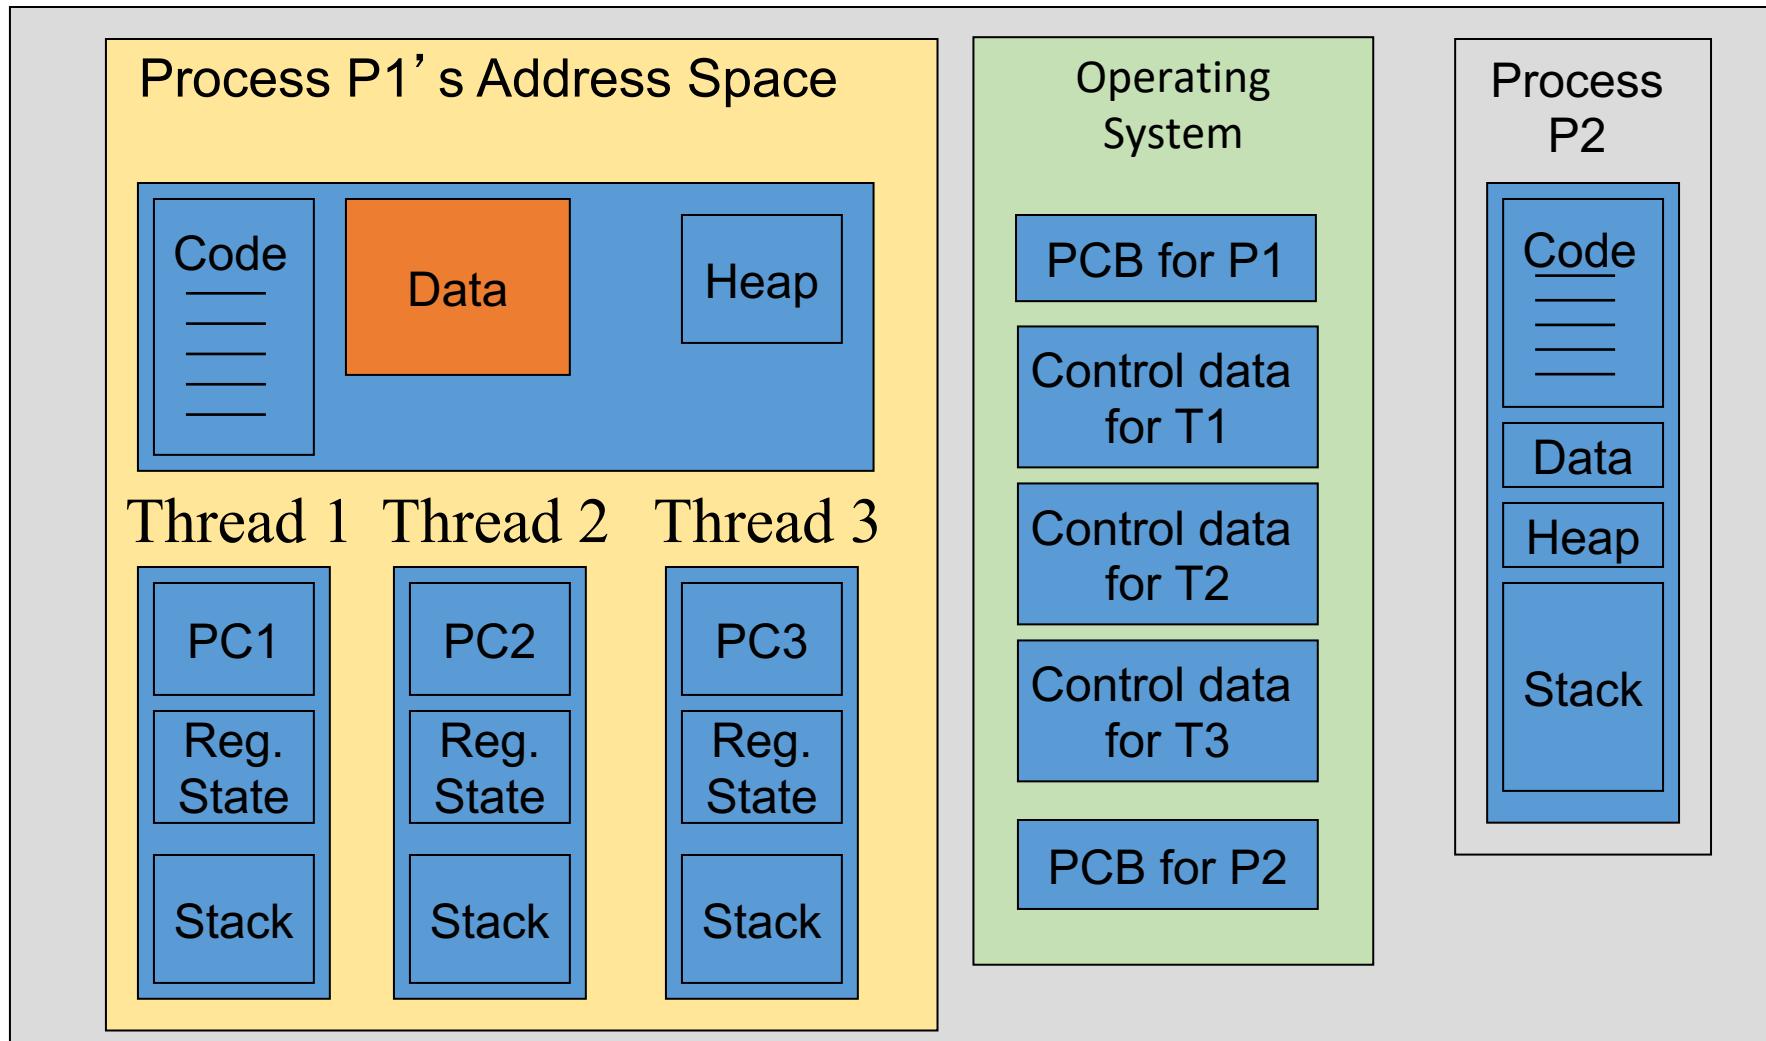




# Threads

# Threads

## Main Memory



# Threads

- Process vs Thread
- Kernel vs User-space threads
  - Many-to-One model
  - One-to-One model
  - Many-to-Many model

# Benefits vs Thread Safety

- Benefits
  - Responsiveness
  - Resource sharing
  - Low context-switching overhead
  - Scalability
- Thread safety
  - Thread safe
    - If the code behaves correctly during simultaneous or concurrent execution by multiple threads
  - Reentrant
    - If the code behaves correctly when a single thread is interrupted in the middle of executing the code reenters the same code

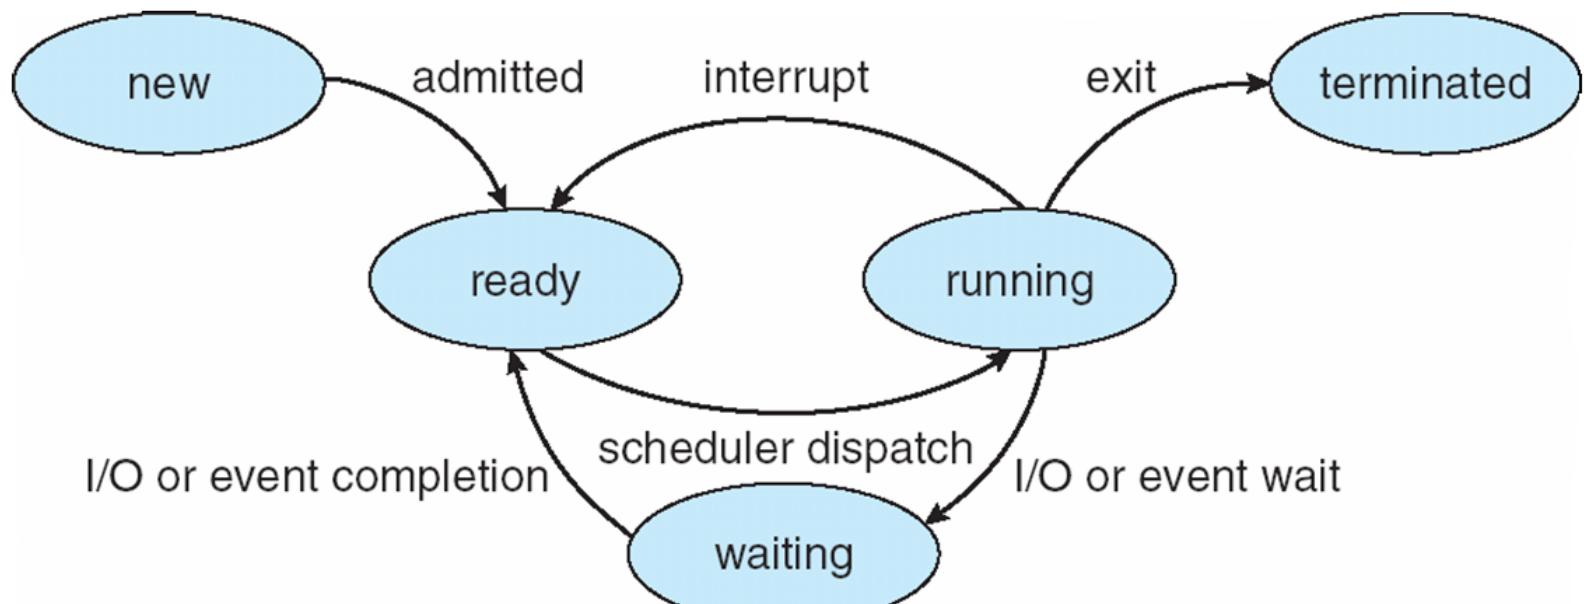
# Synchronization

- Problems without synchronization
  - Race condition
    - Critical section
  - Deadlock
  - Starvation
- Synchronization mechanisms
  - Mutual exclusion
    - Atomic Test & Set
  - Semaphore
  - Monitor
  - Conditional variables

# Synchronization

- Producer – Consumer (Bounded Buffer) problem
  - Problem description?
  - Solution?
- Reader – Writer problem
  - Problem description?
  - Solution?
- Dining Philosophers problem
  - Problem description?
  - Solution?

# Diagram of Process State



Also called “blocked” state



# Deadlock

# Conditions

The following 4 conditions must hold simultaneously for deadlock to arise:

- Mutual exclusion
- Hold and wait
- No preemption
- Circular wait

# Deadlock

- Detection
- Prevention
- Avoidance

**Question 1**

Not yet answered

Points out of 1.00

Below are two processes using the semaphores initialized as Q=1 and S=1. Which of the following is true regarding the execution of these processes?

$P_0$

```
wait(S);      wait(Q);
wait(Q);      wait(S);
.
.
.
signal(S);    signal(Q);
signal(Q);    signal(S);
```

$P_1$

Select one or more:

- a.  $P_1$  is subject to starvation
- b.  $P_0$  and  $P_1$  will always complete their execution
- c.  $P_1$  will always run before  $P_0$
- d.  $P_0$  will always run before  $P_1$
- e.  $P_0$  and  $P_1$  can result in deadlock
- f.  $P_0$  is subject to starvation

**Question 1**

Not yet answered

Points out of 5.00

```
int temp;

void swap(int *y, int *z)
{
    int local;
    local = temp;
    temp = *y;
    *y = *z;
    *z = temp;
    temp = local;
}
```

Select the best answer regarding the given code snippet.

Select one:

- a. The code is re-entrant and thread safe
- b. The code is thread safe but not re-entrant
- c. The code is neither thread safe nor re-entrant
- d. The code is re-entrant but not thread safe

**Question 1**

Not yet answered

Points out of 5.00

```
int temp;

void swap(int *y, int *z)
{
    int local;
    local = temp;
    temp = *y;
    *y = *z;
    *z = temp;
    temp = local;
}
```

Select the best answer regarding the given code snippet.

Select one:

- a. The code is re-entrant and thread safe
- b. The code is thread safe but not re-entrant
- c. The code is neither thread safe nor re-entrant
- d. The code is re-entrant but not thread safe

# Sample Question

**Multiple Choice Questions:** Choose one option that answers the question best.

1. Advantages of threads over processes include
  - A. lower context switch time
  - B. no possibility of race conditions
  - C. sharing of heap and stack
  - D. smaller code size
  - E. All of the above

# Sample Question

**Multiple Choice Questions:** Choose one option that answers the question best.

2. Which of the following is FALSE about IPC via pipes?
  - A. Basic primitives are `send()` and `receive()`.
  - B. Communication can be blocking or non-blocking.
  - C. Pipes can be anonymous or named.
  - D. Pipes can be used for only one-way communication.
  - E. IPC via pipes is slower than IPC using shared memory.

# Sample Question

**Multiple Choice Questions:** Choose one option that answers the question best.

3. Which of the following is NOT required for a system to be in deadlock?
  - A. mutual exclusion
  - B. no preemption
  - C. acquire and hold
  - D. circular dependency

# Sample Question

**Short Answer Questions:** Consider the following program code for the next question (Question 1).

```
int ret = fork();
if (ret == 0){
    ret = fork();
    printf("Hello \n");
}
printf("World \n");
return 0;
```

1. How many times the following word will be printed?

A. “Hello”: \_\_\_\_\_ times

B. “World”: \_\_\_\_\_ times

# Sample Question

**Short Answer Questions:** Consider the following program code for the next question (Question 1).

```
int ret = fork();
if (ret == 0){
    ret = fork();
    printf("Hello \n");
}
printf("World \n");
return 0;
```

1. How many times the following word will be printed?

A. “Hello”: 2 times

B. “World”: 3 times

# Sample Question

**Short Answer Questions:** Consider the following program code for the next question (Question 2).

```
int t;

void swap(int *x, int *y)
{
    int s;
    s = t;
    t = *x;
    *x = *y;
    *y = t;
    t = s;
}
```

2. Is there a possibility of a race condition updating variable t in the code above?

# Sample Question

**Short Answer Questions:** Consider the following program code for the next question (Question 2).

```
int t;

void swap(int *x, int *y)
{
    int s;
    s = t;
    t = *x;
    *x = *y;
    *y = t;
    t = s;
}
```

2. Is there a possibility of a race condition updating variable t in the code above? **YES**

# Sample Question

## Short Answer Questions:

3. Mark each term with the letter of the correct statement.

- |   |  |
|---|--|
| A. A piece of code functions correctly during simultaneous or concurrent execution by multiple threads. | C. A function defined in the Linux kernel to copy data from kernel-space.    |
| B. A function defined in the Linux kernel to copy data to kernel-space.                                 | D. The procedure for replacing the currently executing process with another. |

**Thread-safe**

**copy\_to\_user()**

# Sample Question

## Short Answer Questions:

3. Mark each term with the letter of the correct statement.

- |   |  |
|---|--|
| A. A piece of code functions correctly during simultaneous or concurrent execution by multiple threads. | C. A function defined in the Linux kernel to copy data from kernel-space.    |
| B. A function defined in the Linux kernel to copy data to kernel-space.                                 | D. The procedure for replacing the currently executing process with another. |

**Thread-safe**

A

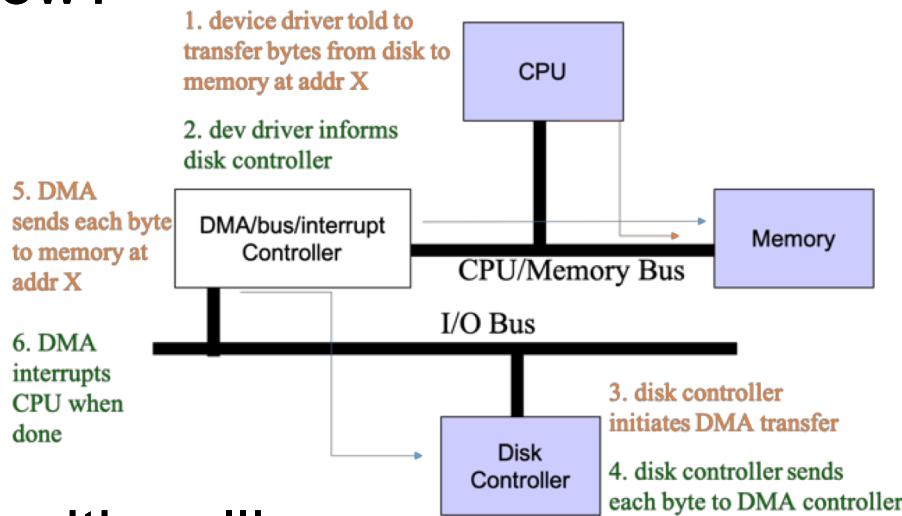
**copy\_to\_user()**

C

# Sample Question

## Short Answer Questions:

4. What is the I/O strategy of the device manager illustrated in the image below?

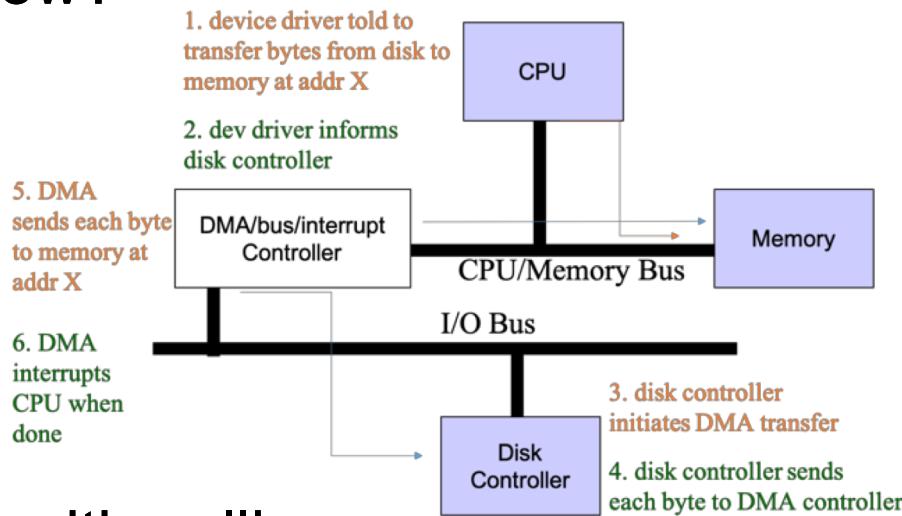


- A. direct I/O with polling
- B. direct I/O with interrupts
- C. DMA with interrupts
- D. hardware interrupts

# Sample Question

## Short Answer Questions:

4. What is the I/O strategy of the device manager illustrated in the image below?



- A. direct I/O with polling
- B. direct I/O with interrupts
- C. DMA with interrupts
- D. hardware interrupts

# Sample Question

1. Consider three processes, P, Q and R with the following code:

- P: ps1; ps2; ps3; ps4;
- Q: qs1; qs2; qs3; qs4;
- R: rs1; rs2; rs3; rs4;

These processes have the following synchronization constraint:

- a) rs1 must be the first statement to execute
- b) process R should be the last process to exit
- c) ps3 must execute after qs2
- d) qs3 must execute after ps4 and rs3
- e) rs4 must execute after either ps4 or qs4 (or both) have executed

Using Condition Variables, provide updated codes for P, Q and R that satisfy these constraints.

# Sample Question

<u>Process P</u>	<u>Process Q</u>	<u>Process R</u>
<ps1>	<qs1>	<rs1>
<ps2>	<qs2>	<rs2>
<ps3>	<qs3>	<rs3>
<ps4>	<qs4>	<rs4>

# Sample Question

<u>Process P</u>	<u>Process Q</u>	<u>Process R</u>
<u>wait(&amp;s1)</u> <ps1>	<u>wait(&amp;s1)</u> <qs1>	<rs1>
<ps2>	<qs2>	<u>signal(&amp;s1)</u> <u>signal(&amp;s1)</u> <rs2>
<ps3>	<qs3>	<rs3>
<ps4>	<qs4>	<rs4>

# Sample Question

<u>Process P</u>	<u>Process Q</u>	<u>Process R</u>
<u>wait(&amp;s1)</u> <ps1>	<u>wait(&amp;s1)</u> <qs1>	<rs1>
<ps2> <u>wait(&amp;s2)</u>	<qs2> <u>signal(&amp;s2)</u>	<u>signal(&amp;s1)</u> <u>signal(&amp;s1)</u> <rs2>
<ps3>	<qs3>	<rs3>
<ps4>	<qs4>	<rs4>

# Sample Question

<u>Process P</u>	<u>Process Q</u>	<u>Process R</u>
<u>wait(&amp;s1)</u> <ps1>	<u>wait(&amp;s1)</u> <qs1>	<rs1>
<ps2> <u>wait(&amp;s2)</u>	<qs2> <u>signal(&amp;s2)</u> <u>wait(&amp;s3)</u>	<rs2>
<ps3>	<qs3>	<rs3>
<ps4> <u>signal(&amp;s3)</u>	<qs4>	<rs4>

# Sample Question

<u>Process P</u>	<u>Process Q</u>	<u>Process R</u>
<u>wait(&amp;s1)</u> <ps1>	<u>wait(&amp;s1)</u> <qs1>	<rs1>
<ps2> <u>wait(&amp;s2)</u>	<qs2> <u>signal(&amp;s2)</u> <u>wait(&amp;s3)</u>	<rs2>
<ps3>	<qs3>	<rs3>
<ps4> <u>signal(&amp;s3)</u> <u>signal(&amp;s4)</u>	<qs4> <u>signal(&amp;s4)</u>	<rs4>

# Sample Question

<u>Process P</u>	<u>Process Q</u>	<u>Process R</u>
<u>wait(&amp;s1)</u> <ps1>	<u>wait(&amp;s1)</u> <qs1>	<rs1>
<ps2> <u>wait(&amp;s2)</u>	<qs2> <u>signal(&amp;s2)</u> <u>wait(&amp;s3)</u>	<rs2>
<ps3>	<qs3>	<rs3>
<ps4> <u>signal(&amp;s3)</u> <u>signal(&amp;s4)</u> <u>signal(&amp;s5)</u>	<qs4> <u>signal(&amp;s4)</u> <u>signal(&amp;s5)</u>	<rs4> <u>wait(&amp;s5)</u>